

Models and Systems for Big Data

DOCUMENT DATABASE - MONGODB

Nacéra Seghouani

Computer Science Department, CentraleSupélec
Laboratoire de Recherche en Informatique, LRI LaHDAK group
nacera.bennacer@centralesupelec

2019-2020

Introduction

- ➡ JSON (JavaScript Object Notation) is used as a data model in NoSQL document-based databases such as MongoDB, CouchDB, CouchBase, RethinkDB
- ➡ JSON: a format to exchange structured data in a distributed environment, used in Ajax Web applications and in Rest architecture based Web services.
- ➡ JSON: simple and intuitive, it hasn't schema specification and query language (generally Javascript)
- ➡ A JSON document is a set of key-value pairs (attribute, value).

📎 Simple values

```
"title": "The Social network",  
"year": 2010,  
"oscar": false
```

📎 Complex values

```
"artist": {"last_name": "Fincher", "first_name": "David"},  
"artists": [{"first_name": "Jesse", "last_name": "Eisenberg"},  
             {"first_name": "Rooney", "last_name": "Mara"}]
```



JSON Data Model



- ✎ Documents are characterized by their structure which can be more or less complex.
- ✎ A document has a tree structure: nodes contain simple values (strings, integers, reals, booleans) or complex values (arrays of simple values or embedded documents). Simple values are the tree leaves.

```
{  
  "title": "The Social network",  
  "summary": "On a fall night in 2003, Harvard undergrad and programming ...",  
  "year": 2010,  
  "director": {"last_name": "Fincher", "first_name": "David"},  
  "actors": [  
    {"first_name": "Jesse", "last_name": "Eisenberg"},  
    {"first_name": "Rooney", "last_name": "Mara"}  
  ]  
}
```

JSON Data Model



- 🔗 Powerful modeling: attributes with no atomic values, embedded/nested documents


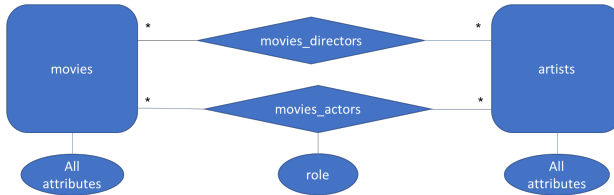
```
[
  {"title": "The Social network",
   "summary": "On a fall night in 2003, Harvard undergrad and programming ...",
   "year": 2010,
   "director": {"last_name": "Fincher", "first_name": "David"},
   "actors": [{ "first_name": "Jesse", "last_name": "Eisenberg"},
               { "first_name": "Rooney", "last_name": "Mara"}] }
  {"title": "Pulp fiction",
   "year": "1994",
   "genre": ["Action", "Crime", "Comedy"]}
  {"country": "USA",
   "director": {"last_name": "Tarantino", "first_name": "Quentin", "birth_date": "1963"},
   "actors": [{ "last_name": "Travolta", "first_name": "John", "role": "Vincent Vega"},
               { "last_name": "Willis", "first_name": "Bruce", "role": "Butch Coolidge" },
               { "last_name": "Jackson", "first_name": "Samuel L.", "role": "Jules Winnfield"},
               ... ]
  }
  ... ]
```



JSON Data Model

- Documents in the same collection are not described using the same attributes
- Hierarchy leads to a non symmetric representation, the access to the root is privileged, and a certain perspective is imposed
- No information scattered among different documents likewise no need to join operations as in relational databases
- No autonomous entities, redundancy and perhaps inconsistencies.


JSON Data Model Versus Graph & Relational Model




Document Database MongoDB

- ☞ MongoDB (from hum**mongous**) is a free and open-source platform database.
- ☞ A powerful, flexible, and scalable NoSQL general-purpose database used in platforms such as: *Craigslist, eBay, Foursquare, SourceForge.net, Viacom, and New York Times.*
- ☞ A document-oriented database
 - ✎ a row is a document, with embedded documents and arrays, complex hierarchical relationships in a single row
 - ✎ a collection is a group of documents, as a table with a dynamic schema
 - ✎ free schema: documents within a collection can have any number of different attributes, not of fixed types or sizes
- ☞ Flexible structure (presence, absence, variations), complex and free imbrication.
- ☞ Few reference to other documents.
- ☞ Transaction in MongoDB . . . and data distribution -> see later

Data Types

 Each document has a unique `_id` within a collection

```
[
  {_id: 1,
  name: "sue",
  age: 19,
  gender : 2,
  favorites: { artist: "Picasso", food: "pizza"},
  badges: [ "blue", "black" ]
}
{_id: 2,
  name: "john",
  age: 21
}
]
```

 Data types:

```
regular expression {"x" : /foobar/i}
array {"x" : ["a", "b", "c"]} {"things" : ["pie", 3.14]}
embedded document {"name":"John Doe","address":
{"street":"123 Park Street","city":"AnyTown","state":"NY"} }
javascript code {"x" : function() { /* ... */ }}
```



Basic Operations

👉 MongoDB comes with a simple but powerful JavaScript shell

- ✎ insert a blog in blog collection (created if it doesn't exist) located in current database
db. _id field is added to each document
- ✎ query the collection

```
> post = {  
  "title": "My Blog Post",  
  "content": "Here's my post",  
  "date" : new Date()  
}  
  
> db.blog.insert(post)  
  
> db.blog.find()  
{ "_id" : ObjectId("5037ee4a1084eb3fffeef7228"),  
  "title" : "My Blog Post",  
  "content" : "Here's my blog post.",  
  "date" : ISODate("2017-04-24T21:12:09.982Z")  
}
```

Basic Operations

- ✎ modify the variable post, add a "comments" attribute

```
> post.comments = []  
  
> db.blog.update({title : "My Blog Post"}, post)  
  
> db.blog.find()  
{  
  "_id" : ObjectId("5037ee4a1084eb3ffeef7228"),  
  "title" : "My Blog Post",  
  "content" : "Here's my blog post.",  
  "date" : ISODate("2012-08-24T21:12:09.982Z"),  
  "comments" : [ ]  
}
```

- ✎ remove documents with the specified attributes values from a collection

```
> db.blog.remove({title : "My Blog Post"})
```

Querying

- ☞ Selection: the first parameter of `find()` specify a set of conditions `{field_i: val_i, field_j.field_k: val_{jk}, ... }` on different attributes as **conjunctive conditions**



```
db.users.find()  
// all documents in users collection
```

```
db.restaurants.find( { "borough": "Manhattan"})  
//documents with attribute "borough" and value "Manhattan"
```

```
db.users.find({"username" : "joe", "age" : 27})  
//documents where "username" : "joe" and "age" : 27
```

```
db.users.find({"username" : "joe", "address.city" : "Paris"})  
db.restaurants.find({"cuisine":"Italian", "address.zipcode":"10075"})  
//condition with embedded document
```

Querying

- 👉 Projection: the second parameter of `find()`
`{field_i: 1, field_j: 0, ... }` 1 if projected, 0 otherwise.
- 👉 `_id` key is returned by default.

```
db.users.find({}, {"username" : 1, "email" : 1})
```

```
db.users.find({}, {"username" : 1, "email" : 1, "_id":0})
```

```
db.restaurants.find({"cuisine":"Italian"}{"address.zipcode": 1})
```

Querying

➡ Selection operators on a single attribute $\{field_i : \{\$operator : val_i\}, \dots\}$

➡ $\$ne$ can be used with any type

```
db.users.find({"age":{"$gte" : 18, "$lte" : 30}})
//users who are between 18 and 30
```

```
start = new Date("01/01/2007")
```

```
db.users.find({"registered" : {"$lt" : start}})
```

```
db.users.find({"username" : ["$ne" : "joe"]})
```

```
db.restaurants.find( { "borough": "Manhattan", "grades.score": { $gt: 30 } })
```

```
db.restaurants.find( { "grades.grade": { $in: [ "A", "B" ] } })
```

```
db.raffle.find({"ticket_no" : {"$nin" : [725, 542, 390] } })
```


Querying

- Selection operators on multiple attributes

$\{ \$operator : [condition_i, condition_j, \dots] \}$

- Two ways to do an `or` query

-  `$in` used to query for a variety of values for a single field.

-  `$or` used to query for any of the given values across multiple fields.

```
db.raffle.find({"$and" : [{"ticket_no" : 725}, {"winner" : true}] })
```

```
db.restaurants.find({$or:[{"cuisine": "Italian"}, {"address.zipcode": "10075"}] })
```

Querying arrays

- 🔗 Querying for an array, for an element i of an array.

```
{field: [val_1, val_2, ...] }; {field: val}; {field.i: val}
```

- 🔗 Querying for arrays of a given size $\$size$

```
db.food.insert({"fruit" : ["apple", "banana", "peach"]})
```

```
db.food.find({"fruit" : "banana"})//at least 1 of array elements matches
```

```
db.food.find({"fruit":["apple","banana","peach"]})  
//exact array match (including order)
```

```
db.food.find({$and : [{fruit:"apple"}, {fruit:"banana"}]})  
//at least 2 of array elements are "apple" and "banana"
```

```
db.food.find({"fruit.2" : "peach"}) //element with index 2 is "peach"
```

```
db.food.find({"fruit" : {"$size" : 3}})
```



Querying arrays

- 👉 Querying for arrays that contain at least one element that matches all the specified query conditions (useful if more than one condition).

```
{field: {$elemMatch : {condition1},{condition2}, ...} }
```

```
db.scores.find({ results: { $elemMatch: { $gte: 80, $lt: 85 } } })
```

- 👉 `$elemMatch` won't match non-array elements.

```
>db.test.find()  
{ "x" : 5 } { "x" : 15 } { "x" : 25 } { "x" : [ 5, 25 ] }  
>db.test.find({ "x" : { "$gt" : 10, "$lt" : 20 } })  
{ "x" : 15 } { "x" : [ 5, 25 ] }  
>db.test.find({ { "x" : { "$elemMatch" : { "$gt" : 10, "$lt" : 20 } } } })  
// no results
```


Querying Arrays

📁 Embedded documents in Arrays

```
>db.blog.find()  
{  
  "content" : "...",  
  "comments" : [  
    {"author" : "joe", "score" : 3, "comment" : "nice post"},  
    {"author" : "mary", "score" : 6, "comment" : "terrible post"}  
  ]  
}
```

```
>db.blog.find({"comments" : {"author" : "joe", "score" : {"$gte" : 5}}})  
// no embedded document
```

```
>db.blog.find({"comments.author": "joe", "comments.score": {"$gte" : 5}})  
// two embedded documents
```

```
>db.blog.find({"comments":{"$elemMatch":{"author":"joe","score":{"$gte" : 5}}}})  
// no embedded document
```

Querying

- ✎ `$slice` operator can be used to return a subset of elements for an array attribute, also by taking an offset and the number of elements to return

```
db.blog.posts.find({}, {"comments" : {"$slice" : 10}}) \\the first 10
```

```
db.blog.posts.find({}, {"comments" : {"$slice" : -10}}) \\the last 10
```

```
db.blog.posts.find({}, {"comments" : {"$slice" : [23, 10]}})  
//skip the first 23 elements and return the 24th through 33th
```

Cursors

```
> for(i=0; i<100; i++) db.collection.insert({x : i});  
  
> var cursor = db.collection.find();  
> while (cursor.hasNext()) {obj = cursor.next();} // do stuff
```

- 🖋 Client-side cursors control the output of a query, limit the number of results, skip over some results
- 🖋 When *find* is called, the client doesn't query the database immediately (only when the cursor start fetching),
- 🖋 The cursor fetches the first 100 results or first 4 MB of results at once (by default) so that the next calls to *next()* or *hasNext()* will not have to make trips to the server.
- 🖋 After the client has run through the first set of results, it will again ask the server for more results with a *getMore()* request (transparent).

Cursors

```
> for(i=0; i<100; i++) db.collection.insert({x : i});  
  
> var cursor = db.collection.find();  
> while (cursor.hasNext()) {obj = cursor.next();} // do stuff
```

- 🔗 On the server side, a cursor takes up memory and resources. Once a cursor runs out of results or the client sends a message telling it to die, the database can free the resources.
- 🔗 A cursor automatically dies when it finishes iterating through the matching results, or after 10 minutes of inactivity.
- 🔗 Many drivers have implemented a function called `immutable`, which tells the database not to time out the cursor. If you turn off a cursor's timeout, you must iterate through all of its results or kill it to make sure it gets closed.

Cursors

- ✎ The most common query options are limiting the number of results returned, skipping a number of results, and sorting. All these options must be added before a query is sent to the database server.

```
> db.c.find().limit(3);  
//only 3 matching documents are returned  
  
> db.c.find().sort({username : 1, age : -1}).limit(3);  
//1 (ascending) or -1 (descending)  
  
> db.c.find().sort({username : 1, age : -1}).skip(3);  
//skip the first 3 matching documents  
//If there are fewer than 3 documents, no results.  
  
> var page1 = db.foo.find({_id}).limit(100)  
  
> var page2 = db.foo.find({}).skip(100).limit(100)  
  
> var page3 = db.foo.find({}).skip(200).limit(100) ...
```

Aggregation

- 📎 MongoDB provides aggregation capabilities using `aggregate`:

```
{ $group: { _id: groupingFields, field1: {function : exp }, ... } }
```

```
$group : { "_id" : "$day" }
```

```
$group : { "_id" : "$grade" }
```

```
$group : { "_id" : { "state" : "$state", "city" : "$city" } }
```

- 📎 Several aggregation operators as \$sum

```
>db.restaurants.aggregate([ { $group: { "_id": "$borough", "count": {$sum:1} } } ])
```

```
{ "_id" : "Staten Island", "count" : 969 }
```

```
{ "_id" : "Brooklyn", "count" : 6086 }
```

```
{ "_id" : "Manhattan", "count" : 10259 }
```

```
{ "_id" : "Queens", "count" : 5656 }
```

```
{ "_id" : "Bronx", "count" : 2338 }
```

```
{ "_id" : "Missing", "count" : 51 }
```

- 📎 `$match` filters documents so that you can run an aggregation on a subset of documents.

```
>db.restaurants.aggregate([
```

```
{ $match: { "borough": "Queens", "cuisine": "Brazilian" } },
```

```
{ $group: { "_id": "$address.zipcode" , "count": { $sum: 1 } } } ])
```

Aggregation

- 🔗 `$project` and `$sort` aggregated documents, to retrieve the five most prolific authors.

```
>db.articles.aggregate(  
  {$project : {"author" : 1}},  
  {$group : {"_id": "$author", "count": {"$sum" : 1}}},  
  {$sort : {"count" : -1}},  
  {"$limit" : 5})  
//groups the authors by name, increments "count"  
// for each document an author appears in.
```

- 🔗 `$unwind` turns each attribute of an array into a separate document.

```
>db.blog.aggregate(  
  {$project : {"comments" : 1}},  
  {$unwind : "$comments"},  
  {$match : {"comments.author" : "Mark"}})
```