

# Models and Systems for Big Data Management

## RELATIONAL DATABASES

Nacéra Seghouani

Computer Science Department, CentraleSupélec  
Laboratoire de Recherche en Informatique, LRI LaHDAK group  
[nacera.bennacer@centralesupelec.fr](mailto:nacera.bennacer@centralesupelec.fr)

2019-2020

# Relational Theory

- ✓ Introduced by E. F. Codd
- ✓ Let  $R(\mathcal{A})$  a relation,  $A_i$  an attribute  $\in \mathcal{A}$  ( $i \in [1..n]$ ), a set of ordered attributes
- ✓  $A_i$  values  $\in D_i$  are atomic
- ✓ Let  $\mathcal{F}$  a set of functional dependencies  $X \rightarrow Y$  defined on  $\mathcal{A}$ ,  $X, Y \subseteq \mathcal{A}$
- ✓ Let  $\mathcal{T}$  a set of tuples such as  $\{(a_1, a_2, \dots, a_n) \in (D_1 \times D_2 \times \dots \times D_n)\}$

 Examples:

**ex.** *schedule(slot, room, professor, module, group);*


*room, slot  $\rightarrow$  module, group, professor*

*professor, slot  $\rightarrow$  module, group, room*

**ex.** *Beer(name, manufacturer);*

*name  $\rightarrow$  manufacturer*

# Relational Algebra

- ✓ Relational algebra is an algebra where operands are relations.
  - ✓ Relational algebra operators are used as a query language for relations in a database.
- 
- Union, intersection, and difference:  $\cup, \cap, \setminus$ .
  - Selection used for picking certain tuples:  $\sigma$ .
  - Projection used for picking certain attributes:  $\pi$
  - Product (Cartesian) and Join used for composition of relations:  $\times, \bowtie$ .

# Relational Algebra : Selection

✓  $R_2 = \sigma_C(R_1)$

- $C$  is a condition that refers to attributes of  $R_1$ .
- $R_2$  is all tuples of  $R_1$  that satisfy  $C$ .

Relation Sells

bar	beer	price
Joe's	Bud	2.50
Joe's	Miller	2.75
Sue's	Bud	2.50
Sue's	Miller	3.00

Relation  $\sigma_{bar="Joe's"}(Sells)$

bar	beer	price
Joe's	Bud	2.50
Joe's	Miller	2.75



# Relational Algebra: Projection



- ✓  $R_2 = \pi_X(R_1)$ ,  $X$  is a subset of  $R_1$  attribute's set.
  - $R_2$  is constructed by looking at each tuple of  $R_1$ , extracting the attributes in  $X$ , in the specified order, and creating from them a tuple for  $R_2$ .
  - Eliminate duplicate tuples, if any.

Relation Sells

bar	beer	price
Joe's	Bud	2.50
Joe's	Miller	2.75
Sue's	Bud	2.50
Sue's	Miller	3.00

Relation  $\pi_{beer,price}(Sells)$

beer	price
Bud	2.50
Miller	2.75
Miller	3.00

# Relational Algebra: Extended Projection

- ✓ Using the same  $\pi_X$  operator, we allow  $X$  to contain arbitrary expressions involving attributes.
  - For instance arithmetic on number attributes

Relation Prices

beer	price min	price min
Bud	2.50	2.75
Miller	2.75	3.75
Heineken	3.00	4.00

Relation  $\pi_{beer, price\ max - price\ min}(Sells)$

beer	price max — price min
Bud	0.25
Miller	1.00
Heineken	1.00

# Relational Algebra: Cartesian Product

✓  $R_3 = R_1 \times R_2$

- Pair each tuple  $t_1$  of  $R_1$  with each tuple  $t_2$  of  $R_2$
- Concatenation  $t_1 t_2$  is a tuple of  $R_3$ .
- If  $R_1$  and  $R_2$  have attribute  $A$  of the same name, use  $R_1.A$  and  $R_2.A$ .

<i>Sells</i>			<i>Bars</i>	
bar	beer	price	bar	addr
Joe's	Bud	2.50	Joe's	Maple St.
Joe's	Miller	2.75	Sue's	River Rd.
Sue's	Bud	2.50		
Sue's	Miller	3.00		

<i>Sells</i> × <i>Bars</i>				
Sells.bar	Sells.beer	Sells.price	Bars.bar	Bars.addr
Joe's	Bud	2.50	Joe's	Maple St.
Joe's	Bud	2.50	Sue's	River Rd.
Joe's	Miller	2.75	Joe's	Maple St.
Joe's	Miller	2.75	Sue's	River Rd.
Sue's	Bud	2.50	Joe's	Maple St.
Sue's	Bud	2.50	Sue's	River Rd.
Sue's	Miller	3.00	Joe's	Maple St.
Sue's	Miller	3.00	Sue's	River Rd.



# Relational Algebra: Theta-Join



✓  $R_3 = (R_1 \bowtie_C R_2) = \sigma_C(R_1 \times R_2).$

<i>Sells</i>			<i>Bars</i>	
bar	beer	price	name	addr
Joe's	Bud	2.50	Joe's	Maple St.
Joe's	Miller	2.75	Sue's	River Rd.
Sue's	Bud	2.50		
Sue's	Miller	3.00		

<i>Sells</i> $\bowtie_{bar=name}$ <i>Bars</i>				
Sells.bar	Sells.beer	Sells.price	Bars.name	Bars.addr
Joe's	Bud	2.50	Joe's	Maple St.
Joe's	Miller	2.75	Joe's	Maple St.
Sue's	Bud	2.50	Sue's	River Rd.
Sue's	Miller	3.00	Sue's	River Rd.



# Relational Algebra: Natural Join

✓  $R_3 = (R_1 \bowtie R_2)$

- Natural join connects two relations by equating attributes of the same name, and projecting out one copy of each pair of equated attributes.

Sells			Bars	
bar	beer	price	bar	addr
Joe's	Bud	2.50	Joe's	Maple St.
Joe's	Miller	2.75	Sue's	River Rd.
Sue's	Bud	2.50		
Sue's	Miller	3.00		

Sells $\bowtie$ Bars			
bar	beer	price	addr
Joe's	Bud	2.50	Maple St.
Joe's	Miller	2.75	Maple St.
Sue's	Bud	2.50	River Rd.
Sue's	Miller	3.00	River Rd.

- A join is also an inner join as opposed to an outer join

# Relational Algebra: Outer-Joins



- ✓ Outer join: avoids dangling tuples, tuples that do not join with anything.

Left (outer) join  $\bowtie$

Right (outer) join  $\bowtie$

Full (outer) join  $R_1 \bowtie R_2 = (R_1 \bowtie R_2) \cup (R_1 \bowtie R_2)$

Sells			Bars	
bar	beer	price	bar	addr
Joe's	Bud	2.50	Joe's	Maple St.
Joe's	Miller	2.75	Sue's	River Rd.
Sue's	Bud	2.50		
Sue's	Miller	3.00		
Max's	Miller	3.00		

Sells $\bowtie$ Bars			
bar	beer	price	addr
Joe's	Bud	2.50	Maple St.
Joe's	Miller	2.75	Maple St.
Sue's	Bud	2.50	River Rd.
Sue's	Miller	3.00	River Rd.
Max's	Miller	3.00	null

# Relational Algebra: Set Operators




- ✓ Let  $R_1(\mathcal{A}_1)$  and  $R_2(\mathcal{A}_2)$  be two relations such that  $\mathcal{A}_1 = \mathcal{A}_2$   
Union  $\cup$ , Intersection  $\cap$ , Difference  $\setminus$  between two relations  $R_1$  and  $R_2$  are similar to set operators.



S1			S2		
bar	beer	price	bar	beer	price
Joe's	Bud	2.50	Max's	Heinken	4
Joe's	Miller	2.75	Sue's	Bud	2.50
Sue's	Bud	2.50			
Sue's	Miller	3.00			

S1 $\setminus$ S2		
bar	beer	price
Joe's	Bud	2.50
Joe's	Miller	2.75
Sue's	Miller	3.00

# Relational Algebra: Complex Expressions


- ✓ Combine operators with parentheses and precedence rules: Expression trees.
  - ✓ Precedence of relational operators: (i)  $[\sigma, \pi]$ ; (ii)  $[\times, \bowtie]$ ; (iii)  $\cap$ ; (iv)  $\cup, \setminus$ .
  - ✓ Give the algebra tree (or expression) for each query. Unary *rename* operator is used to rename a relation (useful to remove ambiguity).
-  Find the names of beers sold by "Sue's" bar
-  Find the names of all the bars that are either on "Maple St." or sell "Bud" for less than \$3
-  Find the bars that sell the same beers with different prices.

# Relational Algebra: Complex Expression


Sells			Bars	
bar	beer	price	bar	addr
Joe's	Bud	2.50	Joe's	Maple St.
Joe's	Miller	2.75	Sue's	River Rd.
Sue's	Bud	2.50		
Sue's	Miller	3.00		
Max's	Miller	3.00		

 Find the names of beers sold by "Sue's" bar

$$\pi_{beer}(\sigma_{bar='Sue's'}(bars))$$

 Find the names of all the bars that are either on "Maple St." or sell "Bud" for less than \$3

$$(\pi_{bar}((\sigma_{addr='MapleSt.'}(bars))) \cup (\pi_{bar}(\sigma_{price < 3 \text{ and } beer='Bud'}(sells))))$$


 Find the bars that sell the same beers with different prices.  
where  $sells1 = sells2 = \text{rename}(sells)$ .

$$(\pi_{sells1.bar, sells2.bar}(\sigma_{sells1.price \neq sells2.price \text{ and } sells1.beer = sells2.beer}(sells1 \times sells2)))$$

# Functional Dependency

## ✓ Functional Dependency (FD):

Let  $R(\mathcal{A})$ , s. t.  $X \subset \mathcal{A}, Y \subset \mathcal{A}, X \cap Y = \emptyset$ ,  $X$  determine  $Y$ :  $X \rightarrow Y \in \mathcal{F}$  if

  $\forall (a^X, a^Y) \text{ and } (a'^X, a'^Y), a^X = a'^X \rightarrow a^Y = a'^Y$

where  $a^X$  is a sequence of  $X$  values of a tuple  $t \in \mathcal{T}$

 *movies(title, director, date, nationality, budget) title, director  $\rightarrow$  date, nationality, budget*

## ✓ Elementary Functional Dependency (EFD) $X \rightarrow Y$ if :

$$\nexists X_i \subset X, X_i \rightarrow Y,$$

 *employees(first\_name, last\_name, category, grade, salary) category, grade  $\rightarrow$  salary*  
*grade  $\rightarrow$  salary*



## ✓ Relation Key:




If  $X \cup Y = \mathcal{A}$  and  $X \rightarrow Y$  is a EFD, then  $X$  is a key of  $R$

 *(title, director) is a key of movies relation*

# Functional Dependency

- ✓ A relation could have more than one key

 *schedule(slot, room, professor, module, group);*  
*room, slot  $\rightarrow$  professor, module, group*  
*professor, slot  $\rightarrow$  room, module, group*

- ✓ Reasoning based on FDs: Armstrong rules (transitivity, augmentation, reflexivity, ...), transitive closure & minimal coverage

# Normalization Theory




- ✓ Normalization is the process of efficiently organizing data in a database
- ✓ Eliminating redundant data; ensuring data dependencies and consistencies (after updates)
  - if  $grade \rightarrow salary$ , the salary should not be duplicated for each employee with the same grade.
  - the salary should not be updated for each employee with the same grade.
- ✓ Three main normal forms are defined.






# Normalization Theory

Let  $R(\mathcal{A}), \mathcal{T}, X \rightarrow Y \in \mathcal{F}$  and  $X$  is a key,  $R$  is of the:

- 1<sup>st</sup> Normal Form (1NF): a value of an attribute is atomic (one indivisible value). 


 phone attribute only one value, address attribute couldn't be divided into sub-attributes (street name, city, ...)

- 2<sup>nd</sup> Normal Form (2NF): non key attributes depend fully on the key

$$\nexists X' \subset X \text{ and } \nexists Y' \subset Y, X' \rightarrow Y' \in \mathcal{F}$$

 *schedule(slot, room, professor, group, module)*

*slot, room  $\rightarrow$  professor, group, module* 

*professor  $\rightarrow$  module* 

*module depends partially on the key (slot, room)*

- 3<sup>rd</sup> Normal Form (3NF): non-key attributes depend directly on the key

$$R \text{ is 2NF and } \nexists Y' \subset Y, Y'' \subseteq Y, Y' \rightarrow Y'' \in \mathcal{F}$$

 *employees(identifier, first\_name, last\_name, category, grade, salary)*

*identifier  $\rightarrow$  first\_name, last\_name, category, grade, salary*

*grade  $\rightarrow$  salary*

*salary depends in a transitive way on the key identifier.*

# Normalization Theory

- ✓ Solution is to define a relational schema of different relations

 Schema 1: (*slot*, *room*), (*professor*, *slot*) are the relation keys

*schedule*(*slot*, *room*, #*professor*, *group*)

*professor\_module*(*professor*, *module*)



Or

*schedule*(*slot*, #*professor*, *room*, *group*)

*professor\_module*(*professor*, *module*)

 Schema 2

*employees*(*identifier*, *first\_name*, *last\_name*, *category*, #*grade*)

*grade\_salary*(*grade*, *salary*)

- ✓ Underlined attributes denotes a key and # denotes a foreign key which references a key of a relation
- ✓ A foreign key value should refer to an existing key value

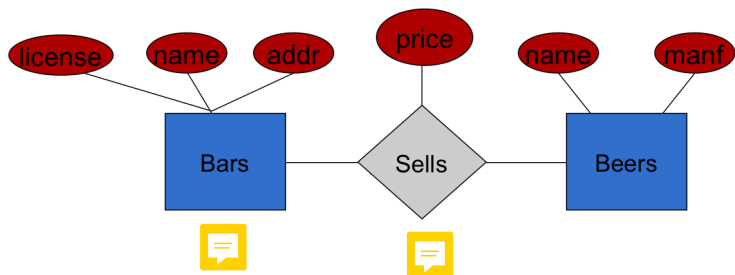
# Normalization Theory



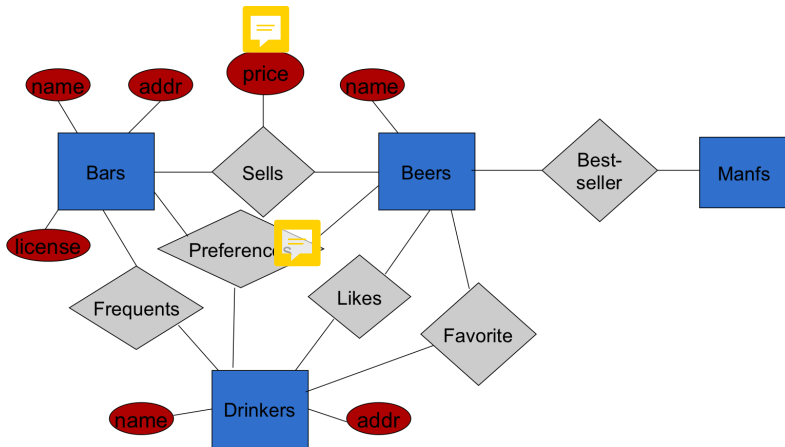
- ✓ Boyce-Codd Normal Form (BCNF):  $R$  is 3NF and  $\forall X \rightarrow Y \in \mathcal{F}$ ,  $X$  is a Key. BCNF takes into account all candidate keys in a relation.
  - Let  $R(A, B, C)$   $A, B \rightarrow C$  and  $C \rightarrow B$   $R$  is in 3NF but not in BCNF because  $C$  is not a key of  $R$
- ✓ Decomposition algorithms to build a schema complying with 1NF, 2NF, 3NF and BCNF

# Entity Association Model

- ✓ Entity defines a collection of similar entities.
- ✓ Attribute is a property of entities of an Entity.
- ✓ Association connects two (binary) or  $n$  (nary) kind of entities.
  - An association can also hold attributes.




# Entity Association Model



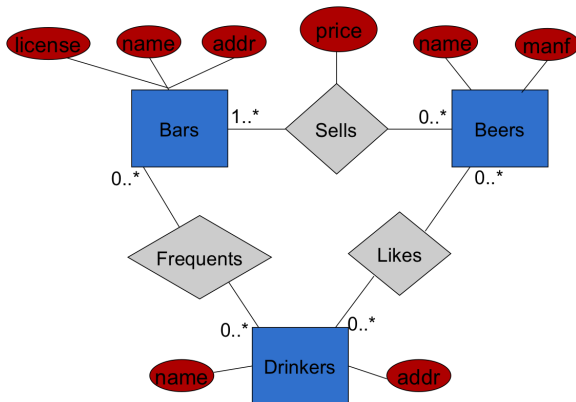
# Entity Association Model

✓ Multiplicity: Associations are Many-Many or Many-One or One-One.



# Mapping EA to Relational Model

- Bars(name, addr, licence); Beers(name, manf); Drinkers(name, addr);  
Sells(#bar, #beer, price); Likes(#drinker, #beer); Frequent(#drinker, #bar)
- The model must at least comply with the functional dependencies and be 2NF and 3NF.



# Relational Algebra - SQL


- ✓ SQL is primarily a query language, for getting information from a relational database. It also includes schema data-definition.  
Relation  $\Rightarrow$  table ; Tuple  $\Rightarrow$  row ; Attribute  $\Rightarrow$  column ; Relational model  $\Rightarrow$  schema
- ✓ A bag is a collection where an element may appear more than once.
- ✓ SQL is a bag language.
  - Bag Laws  $\neq$  Set Laws  
Ex. Set is idempotent for sets  $S \cup S = S$ , but not for bags



## ✓ Creating or deleting a table:

```
CREATE TABLE <name> (<list of declarations> );  
DROP TABLE <name>;
```

## ✓ Declarations include:

 Column's name and its type. The most common types: INT or INTEGER; REAL or FLOAT or NUMERIC(n, nbDecimals); CHAR(n) a fixed-length string; VARCHAR(n) a variable-length string of up to n; DATE or TIME or TIMESTAMP date and time.

 Constraints as Primary Key and Foreign Key

```
CREATE TABLE Beers (name CHAR(20) PRIMARY KEY,  
addr VARCHAR(20), licence CHAR(10));  
CREATE TABLE Bars (name CHAR(20) PRIMARY KEY, manf CHAR(20));  
CREATE TABLE Sells (bar CHAR(20) REFERENCES Bars(name),  
beer CHAR(20) REFERENCES Beers(name), price REAL, PRIMARY KEY (bar, beer));
```

# SQL Language

- ✓ Some other constraints as: UNIQUE, NOT NULL, CHECK

```
CREATE TABLE Persons (  
    ID int UNIQUE,  
    LastName varchar(255) NOT NULL,  
    FirstName varchar(255),  
    Age int CHECK (Age>=18)  
);
```

- ✓ More complex constraints with CHECK and others require procedural language (conditional and repetitive instructions, variables)

# SQL Language

## ✓ SELECT-FROM-WHERE Statements

```
SELECT name FROM Beers WHERE manf = 'Anheuser-Busch';  
SELECT * FROM Beers WHERE manf = 'Anheuser-Busch';  
SELECT name AS beer, manf FROM Beers WHERE manf = 'Anheuser-Busch';  
SELECT bar, beer, price*114 AS priceInYen FROM Sells;  
SELECT price FROM Sells WHERE bar = 'Joe"s Bar' AND beer = 'Bud';  
SELECT price FROM Sells WHERE bar like 'Joe%' AND beer = 'Bud';
```

## ✓ Comparisons:

= , <> , != , > , < , >= , <= , IN , BETWEEN , LIKE , IS NULL , IS NOT NULL

# SQL Language

- ✓ The logic of conditions in SQL is 3-valued logic: TRUE, FALSE, UNKNOWN. 2-valued laws  $\neq$  3-valued laws.

```
NULL = NULL ; NULL <> 1; 1800 + NULL > 1200 -> unknown
```

```
unknown OR false -> unknown
```

```
unknown OR true -> true
```

```
unknown AND false -> false
```

```
unknown AND true -> unknown
```

```
SELECT col FROM t WHERE col = NULL -> unknown
```

```
SELECT col FROM t WHERE col IS NULL -> to use
```



# SQL Language



## Multi-table Queries

```
SELECT beer FROM Likes, Frequenters  
WHERE bar = 'Joe"s Bar' AND Frequenters.drinker = Likes.drinker;  
SELECT b1.name, b2.name FROM Beers b1, Beers b2  
WHERE b1.manf = b2.manf AND b1.name < b2.name;
```





## Join variants

```
R CROSS JOIN S  
R NATURAL JOIN S  
R JOIN S ON <condition>;  
R LEFT/RIGHT/FULL OUTER JOIN S ON <condition>
```



# SQL Language

 Bars(name, addr, licence); Beers (name, manf); Drinkers(name, addr);  
Sells(#bar, #beer, price); Likes(#drinker, #beer); Frequents(#drinker, #bar)

 A parenthesized SELECT-FROM-WHERE statement (sub-queries) can be used as a value in some clauses, including FROM and WHERE clauses.

```
SELECT beer FROM Likes,  
(SELECT drinker FROM Frequents WHERE bar = 'Joe's Bar') JD  
WHERE Likes.drinker = JD.drinker;
```

```
SELECT bar FROM Sells WHERE beer = 'Miller'  
AND price = (SELECT price FROM Sells WHERE bar = 'Joe's Bar' AND beer = 'Bud');
```

```
SELECT * FROM Beers WHERE name IN (SELECT beer  
FROM Likes WHERE drinker = 'Fred');
```

```
SELECT name, man FROM Beers b WHERE NOT EXISTS (SELECT *  
FROM Sells WHERE beer = b.name);
```

EXISTS returns true if at least one result

```
SELECT beer FROM Sells WHERE price >= ALL(SELECT price FROM Sells);
```

ALL returns true if the condition is true for all results

```
SELECT beer, bar FROM Sells WHERE beer = ANY(SELECT beer FROM Likes);
```

ANY returns true if the condition is true for any

# SQL Language

## ✓ Intersection, Union, Difference

```
(SELECT * FROM Likes)
INTERSECT
(SELECT drinker, beer FROM Sells, Frequents WHERE Frequents.bar = Sells.bar);
```

## ✓ Selected columns must be of similar types

## ✓ Duplicates are by default eliminated



## ✓ Force the result to be a bag by ALL as UNION ALL



# SQL Language - Aggregation Operators -

- ✓ Aggregation operators applied to entire column values of a table, produce a single result. Ex. SUM, AVG, COUNT, MIN, MAX.
  - ✎ NULL values are ignored in an aggregation.
  - ✎ COUNT(col) returns the count of non-NULL values of col. COUNT(\*) total number of rows of a table.
  - ✎ If there are no non-NULL values in a col, the result of an aggregation is NULL except COUNT(col) returns 0.
  - ✎ Use DISTINCT inside an aggregation to eliminate duplicates.

```
SELECT AVG(price) FROM Sells WHERE beer = 'Bud';  
SELECT COUNT(DISTINCT price) FROM Sells WHERE beer = 'Bud';  
SELECT bar, MIN(price) FROM Sells WHERE beer = 'Bud' -> illegal
```



# SQL Language - GROUP BY Clause -

- ✓ GROUP BY <list of cols> clause aggregates rows of list columns having the same values and produces a groups for each value.
  - ✎ With GROUP BY, rows with NULL values go into one group, and the aggregates are computed for this group, as for any other.
  - ✎ Projected columns in a GROUP BY must either appear in the GROUP BY clause or under an aggregate function.



```
SELECT beer, AVG(price) FROM Sells GROUP BY beer;
```

```
SELECT AVG(price) FROM Sells GROUP BY bar ;
```

```
SELECT drinker, AVG(price)
FROM Frequents, Sells
WHERE beer = 'Bud' AND Frequents.bar = Sells.bar
GROUP BY drinker;
```

✎ Bars(name, addr, licence); Beers (name, manf); Drinkers(name, addr);  
Sells((#bar, #beer, price); Likes(#drinker, #beer); Frequents(#drinker, #bar)

# SQL Language - HAVING Clause -

- ✓ HAVING <condition> clause may follow a GROUP BY and applies the <condition> to each group, groups not satisfying the <condition> are eliminated.
- ✓ HAVING <condition> without GROUP BY operates on all-at-once the table as a set




```
SELECT beer, AVG(price) FROM Sells GROUP BY beer HAVING AVG(price) > =3  
-> average computed by beer
```

```
SELECT beer, AVG(price) FROM Sells  
GROUP BY beer  
HAVING COUNT(bar) >= 3 OR beer IN (SELECT name FROM Beers WHERE manf = 'Pete''s')
```



```
Bars(name, addr, licence); Beers(name, manf); Drinkers(name, addr);  
Sells(#bar, #beer, price); Likes(#drinker, #beer); Frequents(#drinker, #bar)
```

# SQL Language - INSERT -

 Bars(name, addr, licence); Beers (name, manf); Drinkers(name, addr);  
Sells(#bar, #beer, price); Likes(#drinker, #beer); Frequents(#drinker, #bar)

## Insertion

```
INSERT INTO <relation> VALUES ( <list of values> );
```

```
INSERT INTO Likes VALUES('Sally', 'Bud');
```

```
INSERT INTO Likes(beer, drinker) VALUES('Bud', 'Sally');
```

```
INSERT INTO <relation> ( <subquery> );
```

```
INSERT INTO PotBuddies (SELECT d2.drinker FROM Frequents d1, Frequents d2  
WHERE d1.drinker = 'Sally' AND d2.drinker <> 'Sally' AND d1.bar = d2.bar);
```

# SQL Language - DELETE -



## Deletion



Bars(name, addr, licence); Beers (name, manf); Drinkers(name, addr);  
Sells(#bar, #beer, price); Likes(#drinker, #beer); Frequents(#drinker, #bar)

```
DELETE FROM <relation> WHERE <condition>;
```

```
DELETE FROM Likes WHERE drinker = 'Sally' AND beer = 'Bud';
```

```
DELETE FROM Beers b WHERE EXISTS (SELECT name FROM Beers WHERE manf = b.manf AND  
name <> b.name);
```



```
SELECT b1.name, b2.name FROM Beers b1, Beers b2  
WHERE b1.manf = b2.manf AND b1.name < b2.name;
```

# SQL Language - UPDATE -

## Updates

```
UPDATE <relation> SET <list of attribute assignments> WHERE <condition on tuples>;
```

```
UPDATE Drinkers SET phone = '555-1212' WHERE name = 'Fred';
```

```
UPDATE Sells SET price = 4.00 WHERE price > 4.00;
```